

Optimization on one core
OpenMP, MPI and hybrid programming
*An introduction to the de-facto industrial
standards*

Vincent Keller
Nicolas Richart

July 10, 2019

Table of Contents

Parallelization with OpenMP



Lecture based on specifications ver 3.1

Releases history, present and future

- ▶ October 1997: Fortran version 1.0
- ▶ Late 1998: C/C++ version 1.0
- ▶ June 2000: Fortran version 2.0
- ▶ April 2002: C/C++ version 2.0
- ▶ June 2005: Combined C/C++ and Fortran version 2.5
- ▶ May 2008: Combined C/C++ and Fortran version 3.0
- ▶ **July 2011: Combined C/C++ and Fortran version 3.1**
- ▶ July 2013: Combined C/C++ and Fortran version 4.0
- ▶ November 2015: Combined C/C++ and Fortran version 4.5


Terminology

- ▶ **thread** : an execution entity with a stack and a static memory (*threadprivate memory*)
- ▶ **OpenMP thread** : a *thread* managed by the OpenMP runtime
- ▶ **thread-safe routine** : a routine that can be executed concurrently
- ▶ **processor** : an HW unit on which one or more *OpenMP thread* can execute


Execution and memory models

- ▶ Execution model : fork-join
- ▶ One heavy thread (process) per program (initial thread)
- ▶ lightweight threads for parallel regions. threads are assigned to cores by the OS
- ▶ No implicit synchronization (except at the beginning and at the end of a parallel region)
- ▶ Shared Memory with shared variables
- ▶ Private Memory per thread with threadprivate variables

Memory model (simplified)



Execution model (simplified)



OpenMP and MPI/threads

- ▶ **OpenMP** \neq OpenMPI
- ▶ All what you can do with OpenMP can be done with MPI and/or threads
- ▶ easier **BUT** data coherence/consistency

Syntax in C

OpenMP directives are written as pragmas: `#pragma omp`

Use the conditional compilation flag `#if defined _OPENMP` for the preprocessor

Compilation using the GNU gcc or Intel compiler:

```
gcc -fopenmp ex1.c -o ex1
```

Hello World in C

```
1  #include <stdio.h>
2  #include <omp.h>
3  int main(int argc, char *argv[]) {
4      int myrank=0;
5      int mysize=1;
6      #if defined (_OPENMP)
7      #pragma omp parallel default(shared) private(myrank,
           mysize)
8      {
9          mysize = omp_get_num_threads();
10         myrank = omp_get_thread_num();
11     #endif
12     printf("Hello from thread %d out of %d\n", myrank,
           mysize);
13     #if defined (_OPENMP)
14     }
15     #endif
16     return 0;
```

Syntax in Fortran 90

OpenMP directives are written as comments: !\$omp omp

Sentinels !\$ are authorized for conditional compilation
(preprocessor)

Compilation using the GNU gfortran or Intel ifort compiler:

```
gfortran -fopenmp ex1.f90 -o ex1
```

Number of concurrent threads

The number of threads is specified in a hardcoded way (*omp_set_num_threads()*) or via an environment variable.

BASH-like shells :

```
export OMP_NUM_THREADS=4
```

CSH-like shells :

```
setenv OMP_NUM_THREADS 4
```

Components of OpenMP

- ▶ Compiler directives (written as comments) that allow work sharing, synchronization and data scoping
- ▶ A runtime library (libomp.so) that contains informal, data access and synchronization directives
- ▶ Environment variables

The parallel construct

Syntax

This is the mother of all constructs in OpenMP. It starts a parallel execution.

```
1 #pragma omp parallel [clause[[,] clause]...]
2 {
3     structured-block
4 }
```

where *clause* is one of the following:

- ▶ if or num_threads : conditional clause
- ▶ default(private | firstprivate | shared | none) : default data scoping
- ▶ private(*list*), firstprivate(*list*), shared(*list*) or copyin(*list*) : data scoping
- ▶ reduction({ *operator* / *intrinsic_procedure_name* } : *list*)

Data scoping

What is data scoping ?

- ▶ most common source of errors
- ▶ determine which variables are **private** to a thread, which are **shared** among all the threads
- ▶ In case of a private variable, what is its value when entering the parallel region **firstprivate**, what is its value when leaving the parallel region **lastprivate**
- ▶ The default scope (if none are specified) is **shared**
- ▶ most difficult part of OpenMP

The data sharing-attributes shared and private

Syntax

These attributes determines the scope (visibility) of a single or list of variables

```
1 shared(list1) private(list2)
```

- ▶ The `private` attribute : the data is private to each thread and non-initiatilized. Each thread has its own copy. Example :
`#pragma omp parallel private(i)`
- ▶ The `shared` attribute : the data is shared among all the threads. It is accessible (and non-protected) by all the threads simultaneously. Example :
`#pragma omp parallel shared(array)`

The data sharing-attributes `firstprivate` and `lastprivate`

Syntax

These clauses determines the attributes of the variables within a parallel region:

```
1 firstprivate(list1) lastprivate(list2)
```

- ▶ The `firstprivate` like `private` but initialized to the value before the parallel region
- ▶ The `lastprivate` like `private` but the value is updated after the parallel region

Worksharing constructs

Worksharing constructs are possible in three “flavours” :

- ▶ sections construct
- ▶ single construct
- ▶ workshare construct (only in Fortran)

The single construct

Syntax

```
1 #pragma omp single [clause[[,] clause] ...]  
2 {  
3     structured-block  
4 }
```

where *clause* is one of the following:

- ▶ `private(list), firstprivate(list)`

Only one thread (usually the first entering thread) executes the single region. The others wait for completion, except if the `nowait` clause has been activated

The for directive

Parallelization of the following loop

Syntax

```
1 #pragma omp for [clause[[,] clause] ... ]  
2 {  
3     for-loop  
4 }
```

where *clause* is one of the following:

- ▶ `schedule(kind[, chunk_size])`
- ▶ `collapse(n)`
- ▶ `ordered`
- ▶ `private(list), firstprivate(list),
lastprivate(list),reduction()`

The reduction(...) clause (Exercise)

How to deal with

```
vec = (int*) malloc (size_vec*sizeof(int));
global_sum = 0;
for (i=0;i<size_vec;i++){
    global_sum += vec[i];
}
```

A solution with the reduction(...) clause

```
vec = (int*) malloc (size_vec*sizeof(int));
global_sum = 0;
#pragma omp parallel for reduction(+:global_sum)
    for (i=0;i<size_vec;i++){
        global_sum += vec[i];
    }
```

But other solutions exist !

The schedule clause

Load-balancing

clause	behavior
<i>schedule(static [, chunk_size])</i>	iterations divided in chunks sized <i>chunk_size</i> assigned to threads in a round-robin fashion. If <i>chunk_size</i> not specified system decides.
<i>schedule(dynamic [, chunk_size])</i>	iterations divided in chunks sized <i>chunk_size</i> assigned to threads when they request them until no chunk remains to be distributed. If <i>chunk_size</i> not specified default is 1.

The schedule clause

clause	behavior
<i>schedule(guided [, chunk_size])</i>	iterations divided in chunks sized <i>chunk_size</i> assigned to threads when they request them. Size of chunks is proportional to the remaining unassigned chunks. By default the chunk size is approx $\text{loop_count} / \text{number_of_threads}$.
<i>schedule(auto)</i>	The decisions is delegated to the compiler and/or the runtime system
<i>schedule(runtime)</i>	The decisions is delegated to the runtime system

A parallel for example

How to...

... parallelize the dense matrix multiplication $C = AB$ (triple for loop $C_{ij} = C_{ij} + A_{ik}B_{kj}$). What happens using different schedule clauses ?)

A parallel for example

```
1  #pragma omp parallel shared(A,B,C) private(i,j,k,
    myrank)
2  {
3      myrank=omp_get_thread_num();
4      mysize=omp_get_num_threads();
5      chunk=(N/mysize);
6      #pragma omp for schedule(static, chunk)
7      for (i=0;i<N;i++){
8          for (j=0;j<N;j++){
9              for (k=0;k<N;k++){
10                 C[i][j]=C[i][j] + A[i][k]*B[k][j];
11             }
12         }
13     }
14 }
```

A parallel for example

```
vkeller@mathicsepc13:~$ export OMP_NUM_THREADS=1
```

```
vkeller@mathicsepc13:~$ ./a.out
```

```
[DGEMM] Compute time [s]      : 0.33388209342956
```

```
[DGEMM] Performance  [GF/s]: 0.59901385529736
```

```
[DGEMM] Verification          : 2000000000.00000
```

```
vkeller@mathicsepc13:~$ export OMP_NUM_THREADS=2
```

```
vkeller@mathicsepc13:~$ ./a.out
```

```
[DGEMM] Compute time [s]      : 0.18277192115783
```

```
[DGEMM] Performance  [GF/s]: 1.09425998661625
```

```
[DGEMM] Verification          : 2000000000.00000
```

```
vkeller@mathicsepc13:~$ export OMP_NUM_THREADS=4
```

```
vkeller@mathicsepc13:~$ ./a.out
```

```
[DGEMM] Compute time [s]      : 9.17780399322509E-002
```

```
[DGEMM] Performance  [GF/s]: 2.17917053085506
```

```
[DGEMM] Verification          : 2000000000.00000
```

Synchronization

Synchronization constructs

Those directives are sometimes mandatory:

- ▶ `master` : region is executed by the master thread only
- ▶ `critical` : region is executed by only one thread at a time
- ▶ `barrier` : all threads must reach this directive to continue
- ▶ `taskwait` : all tasks and childs must reach this directive to continue
- ▶ `atomic (read | write | update | capture)` : the associated storage location is accessed by only one thread/task at a time
- ▶ `flush` : this operation makes the thread's temporary view of memory consistent with the shared memory
- ▶ `ordered` : a structured block is executed in the order of the loop iterations

The master construct

- Only the master thread execute the section. It can be used in any OpenMP construct

```
1  #pragma omp parallel default(shared)
2  {
3  ...
4      #pragma omp master
5      {
6          printf("I am the master\n");
7      }
8  ...
9  }
```

Nesting regions

Nesting

It is possible to include parallel regions in a parallel region (i.e. nesting) under restrictions (cf. sec. 2.10, p.111, *OpenMP: Specifications ver. 3.1*)

Runtime Library routines

Usage

- ▶ The functions/subroutines are defined in the lib `libomp.so` / `libgomp.so`. Don't forget to include `#include <omp.h>`
- ▶ These functions can be called anywhere in your programs

Runtime Library routines

Timing routines

routine	behavior
<code>omp_get_wtime</code>	returns elapsed wall clock time in seconds.
<code>omp_get_wtick</code>	returns the precision of the timer used by <code>omp_get_wtime</code>

Environment variables

Usage

- ▶ Environment variables are used to set the ICVs variables
- ▶ under `cs`h : `setenv OMP_VARIABLE "its-value"`
- ▶ under `bash` : `export OMP_VARIABLE="its-value"`

Environment variables

variable	what for ?
OMP_SCHEDULE	sets the run-sched-var ICV that specifies the runtime schedule type and chunk size. It can be set to any of the valid OpenMP schedule types.
OMP_NUM_THREADS	sets the nthreads-var ICV that specifies the number of threads to use in parallel regions

The apparent “easiness” of OpenMP

“Compared to MPI, OpenMP is much easier”

In the reality

- ▶ Parallelization of a non-appropriate algorithm
- ▶ Parallelization of an unoptimized code
- ▶ Race conditions in shared memory environment
- ▶ Memory coherence
- ▶ Compiler implementation of the OpenMP API
- ▶ (Much) more threads/tasks than your machine can support

OpenMP Thread affinity

Affinity = on which core does my thread run ?

Show and set affinity with Intel executable

By setting the export `KMP_AFFINITY=verbose,SCHEDULING` you are able to see where the OS pin each thread

Show and set affinity with GNU executable

By setting the
export `GOMP_CPU_AFFINITY=verbose,SCHEDULING` you are able to see where the OS pin each thread

OpenMP Thread affinity with compact

```
vkeller@mathicsepc13:~$ export KMP_AFFINITY=verbose,compact
vkeller@mathicsepc13:~$ ./ex10
OMP: Info #204: KMP_AFFINITY: decoding x2APIC ids.
OMP: Info #202: KMP_AFFINITY: Affinity capable, using global cpuid leaf 11 info
OMP: Info #154: KMP_AFFINITY: Initial OS proc set respected: {0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15}
OMP: Info #156: KMP_AFFINITY: 16 available OS procs
OMP: Info #157: KMP_AFFINITY: Uniform topology
OMP: Info #179: KMP_AFFINITY: 2 packages x 4 cores/pkg x 2 threads/core (8 total cores)
OMP: Info #206: KMP_AFFINITY: OS proc to physical thread map:
OMP: Info #171: KMP_AFFINITY: OS proc 0 maps to package 0 core 0 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 8 maps to package 0 core 0 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 1 maps to package 0 core 1 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 9 maps to package 0 core 1 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 2 maps to package 0 core 9 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 10 maps to package 0 core 9 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 3 maps to package 0 core 10 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 11 maps to package 0 core 10 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 4 maps to package 1 core 0 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 12 maps to package 1 core 0 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 5 maps to package 1 core 1 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 13 maps to package 1 core 1 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 6 maps to package 1 core 9 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 14 maps to package 1 core 9 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 7 maps to package 1 core 10 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 15 maps to package 1 core 10 thread 1
OMP: Info #144: KMP_AFFINITY: Threads may migrate across 1 innermost levels of machine
OMP: Info #147: KMP_AFFINITY: Internal thread 0 bound to OS proc set {0,8}
OMP: Info #147: KMP_AFFINITY: Internal thread 1 bound to OS proc set {0,8}
OMP: Info #147: KMP_AFFINITY: Internal thread 2 bound to OS proc set {1,9}
OMP: Info #147: KMP_AFFINITY: Internal thread 3 bound to OS proc set {1,9}
[DGEMM] Compute time [s] : 0.344645023345947
[DGEMM] Performance [GF/s]: 0.580307233391397
[DGEMM] Verification : 2000000000.00000
```

OpenMP Thread affinity with scatter

```
vkeller@mathicsepc13:~$ export KMP_AFFINITY=verbose,scatter
vkeller@mathicsepc13:~$ ./ex10
OMP: Info #204: KMP_AFFINITY: decoding x2APIC ids.
OMP: Info #202: KMP_AFFINITY: Affinity capable, using global cpuid leaf 11 info
OMP: Info #154: KMP_AFFINITY: Initial OS proc set respected: {0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15}
OMP: Info #156: KMP_AFFINITY: 16 available OS procs
OMP: Info #157: KMP_AFFINITY: Uniform topology
OMP: Info #179: KMP_AFFINITY: 2 packages x 4 cores/pkg x 2 threads/core (8 total cores)
OMP: Info #206: KMP_AFFINITY: OS proc to physical thread map:
OMP: Info #171: KMP_AFFINITY: OS proc 0 maps to package 0 core 0 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 8 maps to package 0 core 0 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 1 maps to package 0 core 1 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 9 maps to package 0 core 1 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 2 maps to package 0 core 9 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 10 maps to package 0 core 9 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 3 maps to package 0 core 10 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 11 maps to package 0 core 10 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 4 maps to package 1 core 0 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 12 maps to package 1 core 0 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 5 maps to package 1 core 1 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 13 maps to package 1 core 1 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 6 maps to package 1 core 9 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 14 maps to package 1 core 9 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 7 maps to package 1 core 10 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 15 maps to package 1 core 10 thread 1
OMP: Info #144: KMP_AFFINITY: Threads may migrate across 1 innermost levels of machine
OMP: Info #147: KMP_AFFINITY: Internal thread 0 bound to OS proc set {0,8}
OMP: Info #147: KMP_AFFINITY: Internal thread 1 bound to OS proc set {4,12}
OMP: Info #147: KMP_AFFINITY: Internal thread 2 bound to OS proc set {1,9}
OMP: Info #147: KMP_AFFINITY: Internal thread 3 bound to OS proc set {5,13}
[DGEMM] Compute time [s] : 0.204235076904297
[DGEMM] Performance [GF/s]: 0.979263714301724
[DGEMM] Verification : 2000000000.00000
```

OpenMP Thread affinity with explicit (a kind of pinning)

```
vkeller@mathicsepc13:~$ export KMP_AFFINITY='proclist=[0,2,4,6],explicit',verbose
vkeller@mathicsepc13:~$ ./ex10
OMP: Info #204: KMP_AFFINITY: decoding x2APIC ids.
OMP: Info #202: KMP_AFFINITY: Affinity capable, using global cpuid leaf 11 info
OMP: Info #154: KMP_AFFINITY: Initial OS proc set respected: {0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15}
OMP: Info #156: KMP_AFFINITY: 16 available OS procs
OMP: Info #157: KMP_AFFINITY: Uniform topology
OMP: Info #179: KMP_AFFINITY: 2 packages x 4 cores/pkg x 2 threads/core (8 total cores)
OMP: Info #206: KMP_AFFINITY: OS proc to physical thread map:
OMP: Info #171: KMP_AFFINITY: OS proc 0 maps to package 0 core 0 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 8 maps to package 0 core 0 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 1 maps to package 0 core 1 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 9 maps to package 0 core 1 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 2 maps to package 0 core 9 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 10 maps to package 0 core 9 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 3 maps to package 0 core 10 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 11 maps to package 0 core 10 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 4 maps to package 1 core 0 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 12 maps to package 1 core 0 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 5 maps to package 1 core 1 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 13 maps to package 1 core 1 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 6 maps to package 1 core 9 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 14 maps to package 1 core 9 thread 1
OMP: Info #171: KMP_AFFINITY: OS proc 7 maps to package 1 core 10 thread 0
OMP: Info #171: KMP_AFFINITY: OS proc 15 maps to package 1 core 10 thread 1
OMP: Info #144: KMP_AFFINITY: Threads may migrate across 1 innermost levels of machine
OMP: Info #147: KMP_AFFINITY: Internal thread 0 bound to OS proc set {0,8}
OMP: Info #147: KMP_AFFINITY: Internal thread 3 bound to OS proc set {6,14}
OMP: Info #147: KMP_AFFINITY: Internal thread 1 bound to OS proc set {2,10}
OMP: Info #147: KMP_AFFINITY: Internal thread 2 bound to OS proc set {4,12}
[DGEMM] Compute time [s] : 0.248908042907715
[DGEMM] Performance [GF/s]: 0.803509591990774
[DGEMM] Verification : 2000000000.00000
```

“OpenMP-ization” strategy

- ▶ **STEP 1** : Optimize the sequential version:
 - ▶ Choose the best algorithm
 - ▶ “Help the (right) compiler”
 - ▶ Use the existing optimized scientific libraries
- ▶ **STEP 2** : Parallelize it:
 - ▶ Identify the bottlenecks (heavy loops)
 - ▶ “auto-parallelization” is rarely the best !

Goal

Debugging - Profiling - Optimization cycle. Then parallelization !

Tricks and tips

- ▶ **Algorithm:** choose the “best” one
- ▶ **cc-NUMA:** no (real) support from OpenMP side (but OS). A multi-CPU machine is not a real shared memory architecture
- ▶ **False-sharing:** multiple threads write in the same cache line
- ▶ **Avoid barrier.** This is trivial. But sometimes you can't
- ▶ **Small number of tasks.** Try to reduce the number of forked tasks
- ▶ **Asymmetrical problem.** OpenMP is well suited for symmetrical problems, even if tasks can help
- ▶ **Tune the schedule:** types, chunks...
- ▶ **Performance expectations:** a theoretical analysis using the simple Amdahl's law can help
- ▶ **Parallelization level:** coarse (SPMD) or fine (loop) grain ?

What's new with OpenMP 4.0 ?

- ▶ Support for new devices (Intel Phi, GPU,...) with `omp target`. Offloading on those devices.
- ▶ Hardware agnostic
- ▶ League of threads with `omp teams` and distribute a loop over the team with `omp distribute`
- ▶ SIMD support for vectorization `omp simd`
- ▶ Task management enhancements (cancelation of a task, groups of tasks, task-to-task synchro)
- ▶ Set thread affinity with a more standard way than `KMP_AFFINITY` with the concepts of places (a thread, a core, a socket), policies (spread, close, master) and control settings the new clause `proc_bind`